

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import warnings
warnings.filterwarnings('ignore')
```

In [2]:

```
mushroom=pd.read_csv('mushrooms.csv')
```

In [3]:

```
mushroom
#observed that class is the output
```

Out[3]:

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring	stalk-color-above-ring	stalk-color-below-ring	veil-type	veil-color	ring-number
0	p	x	s	n	t	p	f	c	n	k	...	s	w	w	p	w	c
1	e	x	s	y	t	a	f	c	b	k	...	s	w	w	p	w	c
2	e	b	s	w	t	l	f	c	b	n	...	s	w	w	p	w	c
3	p	x	y	w	t	p	f	c	n	n	...	s	w	w	p	w	c
4	e	x	s	g	f	n	f	w	b	k	...	s	w	w	p	w	c
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
8119	e	k	s	n	f	n	a	c	b	y	...	s	o	o	p	o	c
8120	e	x	s	n	f	n	a	c	b	y	...	s	o	o	p	n	c
8121	e	f	s	n	f	n	a	c	b	n	...	s	o	o	p	o	c
8122	p	k	y	n	f	y	f	c	n	b	...	k	w	w	p	w	c
8123	e	x	s	n	f	n	a	c	b	y	...	s	o	o	p	o	c

8124 rows × 23 columns

In [4]:

```
mushroom.describe()
```

Out[4]:

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring	stalk-color-above-ring	stalk-color-below-ring	veil-type	veil-color	ring-number
count	8124	8124	8124	8124	8124	8124	8124	8124	8124	8124	...	8124	8124	8124	8124	8124	8124
unique	2	6	4	10	2	9	2	2	2	12	...	4	9	9	1	4	8
top	e	x	y	n	f	n	f	c	b	b	...	s	w	w	p	w	c
freq	4208	3656	3244	2284	4748	3528	7914	6812	5612	1728	...	4936	4464	4384	8124	7924	7

4 rows × 23 columns

In [5]:

```
mushroom.dtypes
```

Out[5]:

```
class                object
cap-shape            object
cap-surface          object
cap-color            object
bruises             object
odor                object
gill-attachment      object
gill-spacing         object
gill-size            object
gill-color           object
stalk-shape          object
stalk-root           object
stalk-surface-above-ring object
stalk-surface-below-ring object
stalk-color-above-ring object
stalk-color-below-ring object
veil-type           object
veil-color           object
ring-number          object
ring-type            object
spore-print-color     object
population           object
habitat              object
dtype: object
```

In [6]:

```
#converting categorical data in numerical data
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
list1=['class','cap-shape','cap-surface','cap-color','bruises','odor','gill-attachment','gill-spaci
ng','gill-size','gill-color','stalk-shape','stalk-root','stalk-surface-above-ring','stalk-surface-
below-ring','stalk-color-above-ring','stalk-color-below-ring','veil-type','veil-color','ring-number
','ring-type','spore-print-color','population','habitat']
for val in list1:
    mushroom[val]=le.fit_transform(mushroom[val].astype(str))
```

In [7]:

```
mushroom
```

Out[7]:

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring	stalk-color-above-ring	stalk-color-below-ring	veil-type	veil-color	ring-number
0	1	5	2	4	1	6	1	0	1	4	...	2	7	7	0	2	...
1	0	5	2	9	1	0	1	0	0	4	...	2	7	7	0	2	...
2	0	0	2	8	1	3	1	0	0	5	...	2	7	7	0	2	...
3	1	5	3	8	1	6	1	0	1	5	...	2	7	7	0	2	...
4	0	5	2	3	0	5	1	1	0	4	...	2	7	7	0	2	...
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
8119	0	3	2	4	0	5	0	0	0	11	...	2	5	5	0	1	...
8120	0	5	2	4	0	5	0	0	0	11	...	2	5	5	0	0	...
8121	0	2	2	4	0	5	0	0	0	5	...	2	5	5	0	1	...
8122	1	3	3	4	0	8	1	0	1	0	...	1	7	7	0	2	...
8123	0	5	2	4	0	5	0	0	0	11	...	2	5	5	0	1	...


8124 rows × 23 columns

In [8]:

```
# identifying null values
sns.heatmap(mushroom.isnull())
```

Out[8]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1f6a0167248>



In [9]:

```
mushroom.isnull().sum()
```

Out[9]:

```
class 0
cap-shape 0
cap-surface 0
cap-color 0
bruises 0
odor 0
gill-attachment 0
gill-spacing 0
gill-size 0
gill-color 0
stalk-shape 0
stalk-root 0
stalk-surface-above-ring 0
stalk-surface-below-ring 0
stalk-color-above-ring 0
stalk-color-below-ring 0
veil-type 0
veil-color 0
ring-number 0
ring-type 0
spore-print-color 0
population 0
habitat 0
dtype: int64
```

In [10]:

```
mushroom.skew()
```

Out[10]:

```
class 0.071946
cap-shape -0.247052
```

```

cap-shape          0.529951
cap-surface       -0.590859
cap-color          0.706965
bruises            0.342750
odor              -0.080790
gill-attachment   -5.977076
gill-spacing       1.840088
gill-size          0.825797
gill-color         0.061410
stalk-shape       -0.271345
stalk-root         0.947852
stalk-surface-above-ring -1.098739
stalk-surface-below-ring -0.757703
stalk-color-above-ring -1.835434
stalk-color-below-ring -1.791593
veil-type          0.000000
veil-color        -6.946944
ring-number        2.701657
ring-type         -0.290018
spore-print-color  0.548426
population        -1.413096
habitat            0.985548
dtype: float64

```

In [11]:

```
mushroom.corr()
```

Out[11]:

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring	stalk-color-below-ring
class	1.000000	0.052951	0.178446	0.031384	0.501530	0.093552	0.129200	0.348387	0.540024	0.530566	...	0.298801	0.154003
cap-shape	0.052951	1.000000	0.050454	0.048203	0.035374	0.021935	0.078865	0.013196	0.054050	0.006039	...	0.032591	0.031659
cap-surface	0.178446	0.050454	1.000000	0.019402	0.070228	0.045233	-0.034180	0.282306	0.208100	0.161017	...	0.107965	0.066050
cap-color	0.031384	0.048203	0.019402	1.000000	0.000764	0.387121	0.041436	0.144259	0.169464	0.084659	...	0.047710	0.002364
bruises	0.501530	0.035374	0.070228	0.000764	1.000000	0.061825	0.137359	0.299473	0.369596	0.527120	...	0.458983	0.083538
odor	0.093552	0.021935	0.045233	0.387121	0.061825	1.000000	-0.059590	0.063936	0.310495	0.129213	...	0.061820	0.174532
gill-attachment	0.129200	0.078865	0.034180	0.041436	0.137359	0.059590	1.000000	0.071489	0.108984	0.128567	...	0.116177	0.099299
gill-spacing	0.348387	0.013196	0.282306	0.144259	0.299473	0.063936	0.071489	1.000000	0.108333	0.100193	...	0.213775	0.274574
gill-size	0.540024	0.054050	0.208100	0.169464	0.369596	0.310495	0.108984	0.108333	1.000000	0.516736	...	0.010894	0.296548
gill-color	0.530566	0.006039	0.161017	0.084659	0.527120	0.129213	-0.128567	0.100193	0.516736	1.000000	...	0.257224	0.058299
stalk-shape	0.102019	0.063794	0.014123	0.456496	0.099364	0.459766	0.186485	0.080895	0.214576	0.175699	...	0.034399	0.223406
stalk-root	0.379361	0.030191	0.126245	0.321274	0.244188	0.205215	0.144063	0.350548	0.344345	0.315080	...	0.087454	0.157703
stalk-surface-above-ring	0.334593	0.030417	0.089090	0.060837	0.460824	0.118617	-0.088916	0.212359	0.056310	0.224287	...	0.437164	0.132703
stalk-surface-below-ring	0.298801	0.032591	0.107965	0.047710	0.458983	0.061820	-0.116177	0.213775	0.010894	0.257224	...	1.000000	0.106933
stalk-color-above-ring	0.154003	0.031659	0.066050	0.002364	0.083538	0.174532	0.099299	0.274574	0.296548	0.058299	...	0.106933	1.000000
stalk-color-below-ring	0.146730	0.030390	0.068885	0.008057	0.092874	0.169407	0.097160	0.253505	0.278708	0.074781	...	0.110656	0.491503
veil-type	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN

veil-color	0.145142	0.072560	0.016603	0.036130	0.119770	0.057747	0.897518	0.073363	0.103809	0.097583	...	0.072560	0.067518
ring-number	0.214366	0.106534	0.026147	0.005822	0.056788	0.111905	0.093230	0.243014	0.171362	0.096054	...	0.040006	0.084518
ring-type	0.411771	0.025457	0.106407	0.162513	0.692973	0.281387	-0.146689	0.195897	0.460872	0.629398	...	0.394644	0.048618
spore-print-color	0.171961	0.073416	0.230364	0.293523	0.285008	0.469055	-0.029524	0.047323	0.622991	0.416135	...	0.130974	0.271518
population	0.298686	0.063413	0.021555	0.144770	0.088137	0.043623	0.165575	0.529253	0.147682	0.034090	...	0.046797	0.240618
habitat	0.217179	0.042221	0.163887	0.033925	0.075095	0.026610	-0.030304	0.154680	0.161418	0.202972	...	0.039628	0.042518

23 rows × 23 columns



In [12]:

```
for col in mushroom.columns:
    if mushroom.skew().loc[col]>0.55:
        mushroom[col]=np.log1p(mushroom[col])
```

In [13]:

```
#reduced skewness
mushroom.skew()
```

Out[13]:

```
class                0.071946
cap-shape            -0.247052
cap-surface          -0.590859
cap-color            -0.365280
bruises              0.342750
odor                -0.080790
gill-attachment      -5.977076
gill-spacing          1.840088
gill-size             0.825797
gill-color            0.061410
stalk-shape          -0.271345
stalk-root            0.129453
stalk-surface-above-ring -1.098739
stalk-surface-below-ring -0.757703
stalk-color-above-ring -1.835434
stalk-color-below-ring -1.791593
veil-type             0.000000
veil-color           -6.946944
ring-number           1.481287
ring-type            -0.290018
spore-print-color      0.548426
population           -1.413096
habitat               0.342186
dtype: float64
```

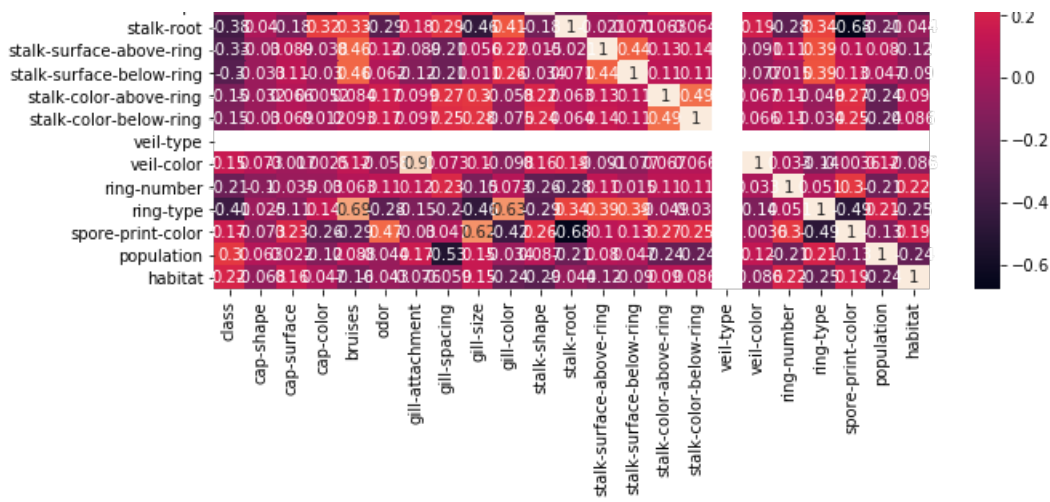
In [14]:

```
plt.figure(figsize=(10,6))
sns.heatmap(mushroom.corr(),annot=True)
```

Out[14]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1f6a09bf4c8>





In [15]:

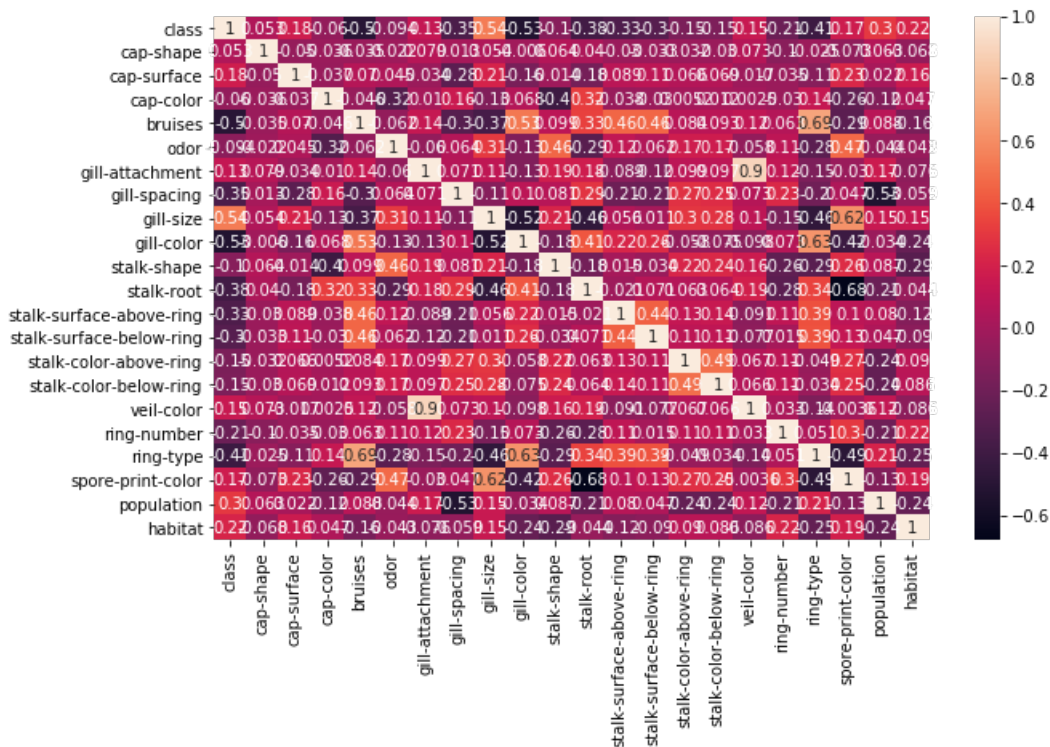
```
mushroom=mushroom.drop(['veil-type'],axis=1)
```

In [16]:

```
plt.figure(figsize=(10,6))
sns.heatmap(mushroom.corr(),annot=True)
```


Out[16]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1f6a1891788>



In [17]:

```
col=mushroom.columns.values
ncol=10
nrow=10
plt.figure(figsize=(ncol,5*ncol))
for i in range(1,len(col)):
    plt.subplot(nrow,ncol,i+1)
    sns.boxplot(mushroom[col[i]],color='green',orient='v')
plt.tight_layout()
```



In [18]:

```
#Removing outliers
from scipy.stats import zscore
z_score=abs(zscore(mushroom))
print(mushroom.shape)
mush=mushroom.loc[(z_score<3).all(axis=1)]
print(mush.shape)
```

(8124, 22)  
(6472, 22)

In [19]:

```
mush
```

Out[19]:

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	...	stalk-surface-above-ring	stalk-surface-below-ring	stalk-color-above-ring	stalk-color-below-ring	c
0	1	5	2	1.609438	1	6	1	0.000000	0.693147	4	...	2	2	7	7	
1	0	5	2	2.302585	1	0	1	0.000000	0.000000	4	...	2	2	7	7	
2	0	0	2	2.197225	1	3	1	0.000000	0.000000	5	...	2	2	7	7	
3	1	5	3	2.197225	1	6	1	0.000000	0.693147	5	...	2	2	7	7	
4	0	5	2	1.386294	0	5	1	0.693147	0.000000	4	...	2	2	7	7	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
8113	1	3	3	1.098612	0	8	1	0.000000	0.693147	0	...	1	1	6	6	
8116	1	3	3	1.609438	0	7	1	0.000000	0.693147	0	...	2	1	6	7	
8117	1	3	2	1.098612	0	8	1	0.000000	0.693147	0	...	1	2	6	7	
8118	1	3	3	1.609438	0	2	1	0.000000	0.693147	0	...	1	2	6	7	
8122	1	3	3	1.609438	0	8	1	0.000000	0.693147	0	...	2	1	7	7	

6472 rows × 22 columns

In [20]:

```
#x and y values allocation for training and testing
x=mush.iloc[:,1:-1]
```

In [21]:

```
x
```

Out[21]:

	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	gill-spacing	gill-size	gill-color	stalk-shape	stalk-root	stalk-surface-above-ring	stalk-surface-below-ring	stalk-color-above-ring	stalk-color-below-ring	s
0	5	2	1.609438	1	6	1	0.000000	0.693147	4	0	1.386294	2	2	7		
1	5	2	2.302585	1	0	1	0.000000	0.000000	4	0	1.098612	2	2	7		
2	0	2	2.197225	1	3	1	0.000000	0.000000	5	0	1.098612	2	2	7		
3	5	3	2.197225	1	6	1	0.000000	0.693147	5	0	1.386294	2	2	7		
4	5	2	1.386294	0	5	1	0.693147	0.000000	4	1	1.386294	2	2	7		
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
8113	3	3	1.098612	0	8	1	0.000000	0.693147	0	1	0.000000	1	1	6		
8116	3	3	1.609438	0	7	1	0.000000	0.693147	0	1	0.000000	2	1	6		
8117	3	2	1.098612	0	8	1	0.000000	0.693147	0	1	0.000000	1	2	6		
8118	3	3	1.609438	0	2	1	0.000000	0.693147	0	1	0.000000	1	2	6		
8122	3	3	1.609438	0	8	1	0.000000	0.693147	0	1	0.000000	2	1	7		

6472 rows × 20 columns

In [22]:

```
x.shape
```

Out[22]:



(6472, 20)

In [23]:

```
y=mush.iloc[:,0]
```

In [24]:

```
y
```

Out[24]:

```
0      1
1      0
2      0
3      1
4      0
```

```
..
8113    1
8116    1
8117    1
8118    1
8122    1
```

Name: class, Length: 6472, dtype: int32

In [25]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=.30,random_state=42)
```

In [26]:

```
#we are using follwing models for prediction
#LogisticRegression
#KNeighborsClassifier
#GaussianNB
#SVC
#DecisionTreeClassifier
#RandomForestClassifier
#AdaBoostClassifier
```

In [27]:

```
lr=LogisticRegression()
lr.fit(x_train,y_train)
lr.score(x_train,y_train)
pred=lr.predict(x_test)
print(accuracy_score(y_test,pred))
print(confusion_matrix(y_test,pred))
print(classification_report(y_test,pred))
```

0.955200823892894

[[990 45]

[ 42 865]]

	precision	recall	f1-score	support
0	0.96	0.96	0.96	1035
1	0.95	0.95	0.95	907
accuracy			0.96	1942
macro avg	0.95	0.96	0.96	1942
weighted avg	0.96	0.96	0.96	1942

In [28]:

```
lrscores=cross_val_score(lr,x,y,cv=5)
print(lrscores)
print(lrscores.mean(),lrscores.std())
```

```
[0.77606178 0.91351351 0.96136012 1.          0.93508501]
0.9172040841901739 0.07623571128762385
```

In [29]:

```
svc=SVC(kernel='rbf')
svc.fit(x_train,y_train)
svc.score(x_train,y_train)
predsvc=svc.predict(x_test)
print(accuracy_score(y_test,predsvc))
print(confusion_matrix(y_test,predsvc))
print(classification_report(y_test,predsvc))
```

```
0.9907312049433573
[[1034    1]
 [   17  890]]
```

	precision	recall	f1-score	support
0	0.98	1.00	0.99	1035
1	1.00	0.98	0.99	907
accuracy			0.99	1942
macro avg	0.99	0.99	0.99	1942
weighted avg	0.99	0.99	0.99	1942

In [30]:

```
svcscores=cross_val_score(svc,x,y,cv=5)
print(svcscores)
print(svcscores.mean(),svcscores.std())
```

```
[0.73050193 0.95057915 0.96445131 1.          0.92581144]
0.9142687664480554 0.09496321361749795
```

In [31]:

```
knn=KNeighborsClassifier()
knn.fit(x_train,y_train)
knn.score(x_train,y_train)
predknn=knn.predict(x_test)
print(accuracy_score(y_test,predknn))
print(confusion_matrix(y_test,predknn))
print(classification_report(y_test,predknn))
```

```
0.9984552008238929
[[1033    2]
 [    1  906]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1035
1	1.00	1.00	1.00	907
accuracy			1.00	1942
macro avg	1.00	1.00	1.00	1942
weighted avg	1.00	1.00	1.00	1942

In [32]:

```
knnscores=cross_val_score(knn,x,y,cv=5)
print(knnscores)
print(knnscores.mean(),knnscores.std())
```

```
[0.76833977 0.98841699 0.9992272 1.          0.96290572]
0.9437779355862818 0.088738502778172
```

In [33]:

```
gnb=GaussianNB()
```

```

gnb=GaussianNB()
gnb.fit(x_train,y_train)
gnb.score(x_train,y_train)
predgnb=gnb.predict(x_test)
print(accuracy_score(y_test,predgnb))
print(confusion_matrix(y_test,predgnb))
print(classification_report(y_test,predgnb))

```

0.8558187435633368

[[985 50]

[230 677]]

	precision	recall	f1-score	support
0	0.81	0.95	0.88	1035
1	0.93	0.75	0.83	907
accuracy			0.86	1942
macro avg	0.87	0.85	0.85	1942
weighted avg	0.87	0.86	0.85	1942

In [34]:

```

gnbscores=cross_val_score(gnb,x,y,cv=5)
print(gnbscores)
print(gnbscores.mean(),gnbscores.std())

```

[0.65559846 0.64633205 0.92194745 1. 0.92581144]

0.8299378778204126 0.14878438201274055

In [35]:

```

dtc=DecisionTreeClassifier()
dtc.fit(x_train,y_train)
dtc.score(x_train,y_train)
preddtc=dtc.predict(x_test)
print(accuracy_score(y_test,preddtc))
print(confusion_matrix(y_test,preddtc))
print(classification_report(y_test,preddtc))

```

1.0

[[1035 0]

[ 0 907]]

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1035
1	1.00	1.00	1.00	907
accuracy			1.00	1942
macro avg	1.00	1.00	1.00	1942
weighted avg	1.00	1.00	1.00	1942

In [36]:

```

dtcscores=cross_val_score(dtc,x,y,cv=5)
print(dtcscores)
print(dtcscores.mean(),dtcscores.std())

```

[0.9011583 1. 1. 1. 0.94435858]

0.9691033758421703 0.0402309391998156

In [37]:

```

rf=RandomForestClassifier()
rf.fit(x_train,y_train)
rf.score(x_train,y_train)
predrf=rf.predict(x_test)
print(accuracy_score(y_test,predrf))
print(confusion_matrix(y_test,predrf))
print(classification_report(y_test,predrf))

```

```

1.0
[[1035    0]
 [    0  907]]
      precision    recall  f1-score   support

         0         1.00      1.00      1.00     1035
         1         1.00      1.00      1.00      907

   accuracy
 macro avg         1.00      1.00      1.00     1942
weighted avg         1.00      1.00      1.00     1942

```

In [38]:

```

rfcores=cross_val_score(rf,x,y,cv=5)
print(rfcores)
print(rfcores.mean(),rfcores.std())

```

```

[0.8023166  1.         1.         1.         0.92581144]
0.9456256079440004 0.07720077199420258

```

In [39]:

```

#DecisionTreeClassifier and RandomForestClassifier are best models among all models
import joblib
joblib.dump(dtc,'mushroom.pkl')

```

Out[39]:

```
['mushroom.pkl']
```

In [ ]: